

EZNX-ATLAS-A: Measuring the Incremental Contribution of Clinical Metadata to 12-Lead ECG Superclass Classification on PTB-XL

Ezyn Segnane^{1,*} , Younes Ommame³ , Khalil El Waled¹ , Mohamedou Cheikh Tourad² , Mariam Mouadili³, Diego Peluffo-Ordóñez⁴  and Mohamed Abdellahi Beddi⁵

¹ Digital Systems and Applications (SAN), Faculty of Sciences and Techniques, University of Nouakchott, Nouakchott, Mauritania; ezyn.segnane@univ-nkc.mr (E.S.); khalil.elwaled@univ-nkc.mr (K.E.W.)

² Al-Khawarizmi, Faculty of Sciences and Techniques, University of Nouakchott, Nouakchott, Mauritania; cheikhtouradmohamedou@univ-nkc.mr (M.C.T.)

³ Faculty of Medical Sciences, UM6P Hospitals, Mohammed VI Polytechnic University (UM6P), 43150 Benguerir, Morocco; Younes.OMMANE@um6p.ma (Y.O.); Maryam.MOUADILI@um6p.ma (M.M.)

⁴ College of Computing, Mohammed VI Polytechnic University (UM6P), Lot 660, Hay Moulay Rachid, 43150 Benguerir, Morocco; PELUFFO.Diego@um6p.ma (D.P.-O.)

⁵ Geometry, Analysis, Algebra and Applications (G3A), Faculty of Sciences and Technology, University of Nouakchott, Nouakchott, Mauritania; mohamed_obeddi@yahoo.fr (M.A.B.)

Abstract

We present a paired multi-seed measurement framework for multimodal-ML ablations, applied to ECG combined with clinical metadata on PTB-XL. We introduce EZNX-ATLAS-A, a 3.95 million-parameter dual-branch architecture coupling a residual ECG encoder, a demographic and anthropometric encoder, an explicit metadata-availability score, and GLU-style late fusion on PTB-XL 1.0.3. We run a **250-run** campaign organised in six experimental groups (A–F) across a total of **270 seed-variant descriptors**. The primary confirmatory group (Group A: three metadata variants $\times$ 20 seeds) reaches macro-AUC 0.9289 (SD 0.0013) for the full-metadata variant. Paired Wilcoxon signed-rank tests with Benjamini–Hochberg correction over the **3-test pre-specified family** (three pairwise macro-AUC contrasts in Group A) reveal that: (i) age and sex alone yield a small but *statistically significant* gain ($\Delta = +0.0007$ macro-AUC, 95% paired CI [0.0002, 0.0011], BH-adjusted $p \approx 0.009$), localised in NORM and MI; (ii) full metadata produce a larger and highly significant gain ($\Delta = +0.0019$ macro-AUC, 95% paired CI [0.0014, 0.0023], BH-adjusted $p < 0.001$), concentrated in NORM, MI, and STTC; and (iii) CD and HYP show no statistically significant effect in either metadata contrast. Exploratory groups confirm that the fusion capacity (Group B, four `meta_hid` values) and the LAUC-BCE balance (Group C, five λ_{LAUC} values) are insensitive to these hyperparameters. Group D demonstrates that training without augmentation incurs a nominally significant macro-AUC drop ($\Delta = -0.0014$, $p = 0.0195$, exploratory). Group E architectural ablations confirm that ECG is essential (meta-only AUC 0.6966, $p = 0.002$) while the gating mechanism variants yield comparable performance. Group F multi-split analysis shows that fold 10 lies below the inter-fold range [0.9400, 0.9445], confirming it is a conservatively difficult evaluation split. Under inference-time masking of anthropometric fields, macro-AUC remains within 0.0014 of the unmasked model, converging to the demographics-only baseline at full masking. These results provide a reproducibility baseline for multimodal ECG studies, demonstrating the value of paired multi-seed statistics, per-class analysis, and explicit missingness audits.

Keywords: electrocardiogram; multimodal fusion; multiplicative gating; missing data; PTB-XL; myocardial infarction; reproducibility; nonparametric inference

Received:

Accepted:

Published:

Copyright: © 2026 by the author.

Submitted to *Mathematics* for possible open access publication under the terms and conditions of the [Creative Commons Attribution \(CC BY\)](https://creativecommons.org/licenses/by/4.0/) license.

MSC: 68T07; 68T45; 62P10; 62G10; 92C55

1. Introduction

Automatic interpretation of the 12-lead electrocardiogram is one of the most mature applications of clinical deep learning. Hannun *et al.* [1] reported cardiologist-level arrhythmia detection on ambulatory single-lead recordings; Ribeiro *et al.* [2] demonstrated large-scale automatic diagnosis from standard 12-lead ECGs; and Attia *et al.* [3] showed that a 12-lead CNN could screen for reduced left-ventricular ejection fraction. Siontis *et al.* [4] review these and related AI-ECG applications in cardiovascular disease management. On the public PTB-XL corpus [5], the canonical benchmark of Strodthoff *et al.* [6] places superclass macro-AUC at 0.928 ± 0.005 for a 1D-ResNet. More recent self-supervised systems [7–9] improve selected ECG benchmarks, but they are waveform-centric and therefore do not directly answer how much structured patient metadata adds to PTB-XL superclass classification.

A natural next step is *multimodal fusion* with structured clinical metadata, since PTB-XL exposes age and sex for nearly all records and height/weight-derived anthropometrics for a smaller subset. Recent ECG-fusion works [10–13] show that ECG representations can be combined with clinical variables, patient-characteristic modules, or engineered feature tables. However, our paper-by-paper audit (Tables 8–9) found that these studies differ in task, dataset, and metadata definition; their reported gains should therefore not be read as direct numerical baselines for PTB-XL superclass demographic/anthropometric fusion. The specific gap we target is narrower: seed-matched ablation of demographics versus anthropometrics, per-class decomposition, and inference-time missingness stress testing.

We contribute a targeted remedy. Rather than proposing yet another ECG backbone, we design **EZNX-ATLAS-A**—a deliberately compact (3.95 M parameters) quality-gated dual-branch architecture—and use it as an *instrument* to address six questions. We intentionally keep a compact residual backbone, rather than escalate simultaneously to a larger Transformer or foundation waveform encoder, so that measured gains are easier to attribute to the metadata pathway itself and so that 250 unique runs across six experimental groups remain computationally tractable:

- Q1. When, if ever, does structured clinical metadata meaningfully improve PTB-XL superclass classification, once seed variance is accounted for?
- Q2. Which *type* of metadata (demographic vs. anthropometric) helps which pathology class?
- Q3. Can a quality-gated fusion degrade gracefully as anthropometric metadata fields are withheld, remaining near the ECG-only and demographics-only baselines?
- Q4. How do our measured effects compare with recent multimodal ECG reports once task, dataset, metadata definition, and statistical protocol are audited?
- Q5. Are the reported gains robust to architectural variants (fusion mechanism, gating, residual path) and is ECG signal necessary?
- Q6. Are the results specific to fold 10, or do they generalise across alternative data splits?

Our answers, detailed in Sections 6–7, are: (Q1) *both* demographic-only and full-metadata variants show statistically significant gains with **20 paired seeds**: demographics alone (demo minus none) yield $\Delta = +0.0007$ macro-AUC (95% paired CI [0.0002, 0.0011],

$p_{\text{BH}} \approx 0.009$), while full metadata (demo+anthro minus none) yield $\Delta = +0.0019$ macro-AUC (95% paired CI [0.0014, 0.0023], $p_{\text{BH}} < 0.001$). (Q2) The class-wise effect of full metadata concentrates in NORM, MI and STTC, while HYP shows no significant gain; (Q3) macro-AUC loss ≈ 0.0014 under full anthropometric masking at inference, converging to the demographics-only baseline; (Q4) the audited literature motivates multimodal fusion but does not provide a directly comparable seed-matched PTB-XL demographic/anthropometric baseline. (Q5) Exploratory Group E ablations (8 architectural variants, 10 seeds each) confirm that the ECG signal is indispensable (meta-only macro-AUC 0.6966, $p = 0.002$) while the specific fusion mechanism is not a driver of performance differences. (Q6) Group F multi-split analysis (4 alternative folds, 5 seeds each) shows that fold 10 lies below the inter-fold range [0.9400, 0.9445], indicating that Group A results are conservatively estimated. These small but reproducible effects support a cautious measurement study replacing broad multimodal claims with statistically honest, class-localised ones.

Beyond the ECG application, this paper introduces a reusable statistical measurement pipeline for multimodal-ML ablations: exact two-sided paired Wilcoxon signed-rank tests on seed-matched differences, Benjamini–Hochberg FDR control at $q = 0.05$ across a **pre-declared 3-test family** (the three pairwise macro-AUC contrasts of Group A), and percentile bootstrap CIs for paired differences. This pipeline applies to any ablation study in which model variants are compared across repeated random seeds.

2. Related Work

2.1. Supervised benchmarks on PTB-XL

Wagner *et al.* [5] released PTB-XL, a 21 799-record 12-lead corpus with SCP-ECG multi-label annotations aggregated into 5 superclasses, 23 sub-classes and 44 diagnostic statements, together with a stratified, patient-disjoint 10-fold partition. Strodtzoff *et al.* [6] established the canonical benchmark, reporting 10-fold cross-validated macro-AUC 0.928 ± 0.005 on super-diag with xresnet1d101. Nonaka and Seita [14] swept nine 1D-CNN backbones and found ResNet-18 competitive with deeper variants. The Bydgoszcz group [15,16] studied entropy features and few-shot learning but used a non-standard 70/15/15 split, making their numbers incomparable to the canonical benchmark. In the supervised PTB-XL benchmark papers audited here, we did not identify a seed-matched age/sex vs. anthropometric ablation with an explicit missing-metadata stress test.

2.2. Landmark clinical ECG models

Hannun *et al.* [1] reached cardiologist-level rhythm classification on 91 232 single-lead ambulatory ECGs; Ribeiro *et al.* [2] generalised to 12-lead short recordings on a 2M-record Brazilian telehealth corpus. Attia *et al.* [3] showed a CNN could screen for LV ejection fraction $\leq 35\%$ directly from the 12-lead ECG (AUC 0.93). Critically for the present paper, Attia *et al.* [17] demonstrated that age and sex are themselves *recoverable* from the raw ECG (age MAE 6.9 y, sex AUC 0.97), supporting the idea that demographic signal is already partly present in the waveform. Siontis *et al.* [4] identify this latent-demographic phenomenon as both an opportunity and a confounding risk.

2.3. Self-supervised and foundation ECG models

Contrastive and masked-modelling approaches—CLOCS [18], CPC-based pretraining [7], ST-MEM [19], the ECG joint-embedding predictive architecture (ECG-JEPA) [8], and preprint waveform foundation models such as HuBERT-ECG [9] and TolerantECG [20]—study representation learning and robustness from waveform inputs alone, including PTB-XL downstream evaluations in some cases. They are therefore adjacent to, but distinct

from, the metadata-value question studied here. The non-peer-reviewed preprints in this paragraph are cited as contextual neighbouring work rather than as pre-specified baselines.

2.4. Multimodal ECG + tabular fusion

Explicit ECG + tabular or feature-fusion architectures have begun to appear, but they are not interchangeable baselines. Qiu *et al.* [10] fuse a transformer ECG encoder with clinical features for atrial-fibrillation recurrence prediction after ablation, outside PTB-XL. The preprint rECGnition_v2.0 system [11] studies self-attentive canonical fusion on MIT-BIH, INCARTDB and the European ST-T Database. Among PTB-XL-family studies, Atwa *et al.* [12] (whose published title uses the abbreviation “PTB-X”, referring to the same PhysioNet corpus) is the closest published demographic-fusion comparator because it adds age and sex to CNN/VGG-style arrhythmia models; however, it reports accuracy-oriented results across 2-, 5-, 10-, and 15-class tasks with augmentation rather than seed-matched macro-AUC on the canonical five-superclass split. In parallel, PTB-XL+ [21] augments PTB-XL with engineered Uni-G, 12SL and ECGDeli feature tables; Holloway *et al.* [13] compare raw PTB-XL signals with such PTB-XL+ feature representations. These feature tables are derived from the ECG itself rather than from EHR-like anthropometrics. We therefore treat these papers as motivation and methodological context, not as direct numerical comparators; the exact comparability audit is summarised in Tables 8–9.

2.5. Gated fusion mechanisms

The Gated Linear Unit (GLU) of Dauphin *et al.* [22] provides a simple multiplicative gate $\mathbf{h} \odot \sigma(\cdot)$ that preserves a direct linear gradient path through the ungated branch. The Gated Multimodal Unit [23] extends this to modality-level mixing. In medical imaging, Chen *et al.* [24] use feature disentanglement plus gated fusion for brain-tumor segmentation, and Kim *et al.* [25] introduce reliability-aware gating for degraded or low-quality modalities. We use a GLU-based gate conditioned on an explicit scalar quality score q_{meta} derived from the metadata-availability mask.

2.6. Missing modalities in clinical ML

Baltrušaitis *et al.* [26] survey the main settings of multimodal learning, including fusion and co-learning under incomplete modalities. HeMIS [27], SMIL [28], and mmFormer [29] address missing modalities in imaging or general multimodal settings rather than ECG+tabular fusion. We study the adjacent ECG+tabular setting by coupling an availability-mask-driven quality score with stochastic sample meta-dropout during training.

2.7. Clinical motivation for anthropometrics in ECG classification

Classical ECG LVH criteria—Sokolow–Lyon [30], Cornell voltage/duration product [31], Romhilt–Estes [32], and Peguero–Lo Presti [33]—depend on measured voltage and are affected by body habitus. Left ventricular mass is itself routinely indexed to body surface area or height^{2.7}. *A priori*, these cardiology facts suggest that the HYP superclass could benefit from anthropometric metadata. More broadly, ECG amplitudes and repolarisation measurements relevant to MI and STTC can also be influenced by torso conduction properties, lead geometry, and recording context [34,35], suggesting possible relevance of body-size priors across morphology-rich classes. Our per-class results in Section 6 are compatible with this broader relevance while not supporting a strong HYP-specific effect at the superclass level.

3. Dataset and Preprocessing

3.1. PTB-XL 1.0.3

We use PTB-XL version 1.0.3 on PhysioNet [5,36] at 100 Hz sampling rate ($T=1000$ samples, 10 s, 12 leads). Following [6], SCP-ECG diagnostic statements are mapped to the 5-class super-diagnostic task $\{\text{NORM}, \text{MI}, \text{STTC}, \text{CD}, \text{HYP}\}$ using `scp_statements.csv`. In the implemented loader, a record is positive for a superclass if any diagnostic SCP code assigned to that class is present in `scp_codes`; no additional likelihood threshold is applied at load time. The official stratified-by-label and patient-disjoint 10-fold partition is used unchanged in the working index: folds 1–8 for training, fold 9 for validation, fold-10 for test, yielding 17,418 / 2,183 / 2,198 records, respectively (Table 4). These fold counts sum to 21 799 records, matching the current PhysioNet v1.0.3 release after the duplicate removals documented in its release notes.

3.2. Formal multi-label problem setting

For each record i , the loader returns a quadruple

$$(\mathbf{X}_i, \mathbf{x}_i, \mathbf{m}_i, \mathbf{y}_i) \in \mathbb{R}^{12 \times 1000} \times \mathbb{R}^8 \times \{0, 1\}^8 \times \{0, 1\}^5,$$

where $\mathbf{X}_i$ is the ECG, $\mathbf{x}_i$ is the imputed metadata vector, $\mathbf{m}_i$ is its availability mask, and $\mathbf{y}_i$ is the multi-hot superclass target. A model with parameters θ emits logits $\ell_i(\theta) \in \mathbb{R}^5$ and probabilities $\mathbf{p}_i = \sigma(\ell_i)$. The principal metric is

$$\text{macroAUC}(\theta) = \frac{1}{5} \sum_{c=1}^5 \text{AUC}(\{(y_{ic}, p_{ic})\}_{i \in \mathcal{D}_{\text{test}}}),$$

with class-wise AUC skipped only if a class is degenerate, which does not occur for the PTB-XL superclasses on fold-10. The ablations below are deterministic interventions on $(\mathbf{x}_i, \mathbf{m}_i)$, so seed-matched comparisons estimate the incremental contribution of metadata while holding the architecture, optimiser, folds, and random seed fixed.

3.3. Metadata extraction and normalisation

From the PTB-XL metadata we construct the 8-dimensional tabular vector

$$\mathbf{x}_{\text{meta}} = \left[\underbrace{\text{age}_z, \text{sex01}}_{\text{demographic}}, \underbrace{\text{height}_z, \text{weight}_z, \text{bmi}_z}_{\text{anthropometric}}, \underbrace{m_h, m_w, m_{\text{bmi}}}_{\text{missingness indicators}} \right] \in \mathbb{R}^8, \quad (1)$$

where $\text{sex01} \in \{0, 1\}$, subscript z denotes standardisation using training-fold mean and standard deviation after median imputation, BMI is derived from height and weight when both are observed and otherwise imputed, and $m \in \{0, 1\}$ is a derived missingness flag set to 1 if the corresponding anthropometric field is absent (0 if originally observed); note that this is opposite to the availability mask $\mathbf{m}$ described below, which uses 1 to indicate presence. $\mathbf{x}_{\text{meta}}$ therefore bundles both imputed feature values and the three derived missingness indicators into the same 8-dimensional vector. The separate availability mask $\mathbf{m} \in \{0, 1\}^8$ (index columns `mask_*`) flags whether each entry of $\mathbf{x}_{\text{meta}}$ was originally present; because m_h, m_w, m_{bmi} are derived binary constants, their mask bits are always 1.

3.4. Positive class weights from training-set prevalence

Let $n_{\text{pos},c}$ and $n_{\text{neg},c}$ denote the numbers of positive and negative training-fold records for superclass c , and let $f_c = n_{\text{pos},c} / (n_{\text{pos},c} + n_{\text{neg},c})$ be its prevalence. We derive the binary cross-entropy positive weights as the clipped negative-to-positive ratio:

$$w_c^+ = \text{clip}\left(\frac{n_{\text{neg},c}}{n_{\text{pos},c}}, 0.5, 30.0\right), \quad c \in \{\text{NORM}, \text{MI}, \text{STTC}, \text{CD}, \text{HYP}\}, \quad (2)$$

Weights grow with class rarity: the most-prevalent class (NORM, 43.6% of training records) receives the smallest weight, while the rarest (HYP, 12.2%) receives the largest. The resulting weight vector $[w_{\text{NORM}}^+, w_{\text{MI}}^+, w_{\text{STTC}}^+, w_{\text{CD}}^+, w_{\text{HYP}}^+] \approx [1.29, 2.98, 3.16, 3.46, 7.22]$, verified from the logged training outputs, is computed dynamically at the start of each run from the training folds only (never from validation or test folds).

3.5. Preprocessing

Signals are loaded with `wfdb` at 100 Hz, padded or truncated to 1000 time samples, transposed to channel-first shape $(12, T)$, cast to `float32`, and divided by 5.0 before entering the model. Before imputation, out-of-range metadata are converted to missing using conservative physiological bounds: age 0–120 years, height 120–210 cm, weight 30–250 kg, and BMI 10–60. Evaluation splits apply only the fixed divisor; training additionally uses light stochastic augmentation. The augmentation block fires with probability 0.5 per training mini-batch; inside it, three operations each trigger independently at probability 0.3: additive Gaussian noise ($\sigma=0.02$), a random integer temporal shift (uniform on $[-20, +19]$ samples), and a multiplicative amplitude jitter (uniform from $[0.95, 1.05]$).

4. Model Architecture: EZNX-ATLAS-A

Figure 1 shows the complete architecture. We describe each block with its dimensional signature and design rationale.




Figure 1. EZNX-ATLAS-A architecture. The 1D residual ECG backbone emits a 1024-dim temporal-statistics embedding $\mathbf{h}_{ts1}$. A dual demographic/anthropometric MLP stack produces $\mathbf{h}_m \in \mathbb{R}^{128}$, scaled by an explicit availability score q_{meta} . One zero-initialised 1024×128 residual projection ($\mathbf{W}_{res}$) softly biases the ECG embedding via residual injection (Eq. 8); a second zero-initialised projection ($\mathbf{W}_{addon}$) is instantiated exclusively for the E3 additive-injection ablation (Group E). A GLU-style multiplicative gate fuses the concatenated embeddings. Three classification heads (fused, ECG-only, meta-only) enable multi-loss training; inference uses the fused-branch outputs ($w^*=1.0$ fixed a priori in all runs). Table 5 reports the matched parameter budget (3.95 M) and the archived compute provenance.

4.1. ECG branch: TSBBackbone1D_v5

An initial Conv1D ($k=7, s=2, p=3, 12 \rightarrow 64$) with BatchNorm and ReLU is followed by MaxPool ($k=3, s=2, p=1$). Three residual stages, each with two residual blocks, implement the channel schedule $64 \rightarrow 64$ (stride 1), $64 \rightarrow 128$ (stride 2), $128 \rightarrow 256$ (stride 2), giving an

output of shape $256 \times T'$ with $T'=63$ for $T=1000$. A *Temporal Statistics Pool* concatenates mean, standard deviation, max, and min over time:

$$\mathbf{h}_{ts1} = [\mu(\mathbf{F}), \sigma(\mathbf{F}), \max(\mathbf{F}), \min(\mathbf{F})] \in \mathbb{R}^{1024}, \quad \mathbf{F} \in \mathbb{R}^{256 \times T'}. \quad (3)$$

4.2. Dual metadata encoder

The loader returns the 8-dim tabular vector in Eq. (1) together with its 8-dim availability mask $\mathbf{m}$. Each branch receives its feature slice concatenated with the corresponding mask slice, yielding a 4-dim demographic input (2 values + 2 availability bits) and a 12-dim anthropometric input (6 values + 6 availability bits):

$$\mathbf{h}_{\text{demo}} = \text{MLP}_{\text{demo}}(\mathbf{x}_{\text{demo}}) : \mathbb{R}^4 \rightarrow \mathbb{R}^{64} \rightarrow \mathbb{R}^{64} \rightarrow \mathbb{R}^{64}, \quad (4)$$

$$\mathbf{h}_{\text{anthro}} = \text{MLP}_{\text{anthro}}(\mathbf{x}_{\text{anthro}}) : \mathbb{R}^{12} \rightarrow \mathbb{R}^{96} \rightarrow \mathbb{R}^{96} \rightarrow \mathbb{R}^{64}. \quad (5)$$

Each intermediate linear layer uses ReLU, BatchNorm, and 10% Dropout; the final linear layer uses ReLU only.

4.3. Availability-aware quality score

We define per-block availability scores and an aggregate metadata-quality score:

$$q_d = \frac{1}{2} \sum_{i=0}^1 m_i, \quad q_a = \frac{1}{6} \sum_{i=2}^7 m_i, \quad q_{\text{meta}} = \min(1, q_d + 0.5 q_a). \quad (6)$$

In the working index, sex is always observed and age is observed for 98.7% of records, so q_d is usually 1, can drop to 0.5 when age is missing, and reaches 0 only in the theoretical edge case where both are unavailable (not observed in PTB-XL). Consequently, $q_{\text{meta}} = \min(1, q_d + 0.5 q_a) = 1$ for virtually all records in the demo and demo+anthro conditions (since $q_d=1$ implies $q_d + 0.5 q_a \geq 1$). Graceful degradation under anthropometric missingness is therefore driven primarily by the explicit missingness flags (m_h, m_w, m_{bmi}) in $\mathbf{x}_{\text{meta}}$ and by sample meta-dropout; q_{meta} scaling provides the hard gate between the none variant ($q_{\text{meta}}=0$) and any variant with demographics present ($q_{\text{meta}}=1$).

4.4. Meta-fusion MLP with sample meta-dropout

Concatenation of the encoder outputs with the two availability scalars produces:

$$\begin{aligned} \tilde{\mathbf{h}}_{\text{anthro}} &= q_a \mathbf{h}_{\text{anthro}}, \\ \mathbf{u}_m &= [\mathbf{h}_{\text{demo}} \parallel \tilde{\mathbf{h}}_{\text{anthro}} \parallel q_d \parallel q_a] \in \mathbb{R}^{130}, \\ \mathbf{h}_m &= q_{\text{meta}} \text{MLP}_m(\mathbf{u}_m) \in \mathbb{R}^{128}. \end{aligned} \quad (7)$$

During training with probability $p=0.10$, the entire pre-quality embedding is zeroed for the sample (*sample meta-dropout*); at inference the dropout is disabled and q_{meta} scaling is applied.

4.5. Residual injection

$$\mathbf{h}_{ts} = \mathbf{h}_{ts1} + 0.10 q_{\text{meta}} \mathbf{W}_{\text{res}} \mathbf{h}_m, \quad \mathbf{W}_{\text{res}} \in \mathbb{R}^{1024 \times 128}, \mathbf{W}_{\text{res}} \leftarrow \mathbf{0} \text{ at init.} \quad (8)$$

The projection $\mathbf{W}_{\text{res}}$ is zero-initialised, making the injection path neutral at initialisation. A second independent zero-initialised projection $\mathbf{W}_{\text{addon}} \in \mathbb{R}^{1024 \times 128}$ is instantiated exclusively for the E3 additive-injection ablation (Group E; Section 6.7) and is inert in the standard architecture (it receives zero gradient and remains at zero throughout standard-

mode training). Because $\mathbf{h}_m$ is itself scaled by q_{meta} (Eq. (7)), the effective attenuation of the residual path is q_{meta}^2 , aggressively suppressing metadata-to-ECG cross-talk at partial availability.

4.6. GLU-style fusion gate

$$\mathbf{h} = [\mathbf{h}_{\text{ts}} \parallel \mathbf{h}_m] \in \mathbb{R}^{1152}, \quad \mathbf{z} = \mathbf{h} \odot \sigma(\text{Linear}_2(\text{ReLU}(\text{Linear}_1(\mathbf{h})))) \in \mathbb{R}^{1152}. \quad (9)$$

4.7. Three heads with meta-correction

$$\ell_{\text{ecg}} = \mathbf{W}_{\text{ecg}} \mathbf{h}_{\text{ts1}} + \mathbf{b}_{\text{ecg}}, \quad (10)$$

$$\ell_{\text{meta}} = \mathbf{W}_{\text{meta}} \mathbf{h}_m + \mathbf{b}_{\text{meta}}, \quad (11)$$

$$\ell_{\text{fused}} = \mathbf{W}_{\text{fused}} \mathbf{z} + \mathbf{b}_{\text{fused}} + 0.05 q_{\text{meta}} \ell_{\text{meta}}. \quad (12)$$

ℓ_{ecg} is computed on $\mathbf{h}_{\text{ts1}}$ before the residual injection in Eq. (8), preserving a strictly unimodal ECG branch. The $0.05 q_{\text{meta}}$ term in Eq. (12) provides a small direct metadata-logit correction to the fused head, gated by availability; at $q_{\text{meta}}=0$ the fused head reduces to $\mathbf{W}_{\text{fused}} \mathbf{z} + \mathbf{b}_{\text{fused}}$ alone.

4.8. Parameter budget

The total is 3.95 M trainable parameters: 0.96 M in the ECG backbone, 0.06 M in the dual metadata encoders and metadata-fusion MLP, 2.66 M in the GLU fusion gate, 0.26 M in the two metadata-to-ECG residual projections ($\mathbf{W}_{\text{res}}$, active in the standard architecture, and $\mathbf{W}_{\text{addon}}$, active in the E3 additive-injection ablation only; each $\mathbb{R}^{1024 \times 128}$, both zero-initialised), and 11.5 k in the three classification heads. The GLU gate accounts for about 67% of trainable parameters.

5. Training Protocol and Experimental Design

5.1. Composite multi-task loss

We minimise

$$\begin{aligned} \mathcal{L} = & 0.52 \text{BCE}_{\mathbf{w}^+}(\ell_{\text{fused}}, \mathbf{y}) + 0.30 \text{BCE}_{\mathbf{w}^+}(\ell_{\text{ecg}}, \mathbf{y}) \\ & + 0.10 q_{\text{meta}} \text{BCE}(\ell_{\text{meta}}, \mathbf{y}) + 0.08 \mathcal{L}_{\text{AUC}}(\sigma(\ell_{\text{fused}}), \mathbf{y}), \end{aligned} \quad (13)$$

where $\text{BCE}_{\mathbf{w}^+}$ is weighted multi-label binary cross-entropy with class weights from Eq. (2) and $\mathcal{L}_{\text{AUC}}$ is an unsquared hinge AUC surrogate pooling all positive/negative label positions across the mini-batch.

5.2. Optimiser and schedule

AdamW with learning rate 10^{-3} , weight decay 5×10^{-4} ; CosineAnnealingWarmRestarts scheduler with $T_0=10$, $T_{\text{mult}}=2$ (with $T_0=10$ matching the epoch budget, the schedule executes one complete cosine cycle and no warm restart is triggered); batch size 32, gradient accumulation 2 (effective batch 64); gradient clipping at $\|\nabla\|_2 \leq 1.0$; 10 epochs. The saved checkpoint is the epoch with the highest validation macro-AUC of the fused branch (with $w^*=1.0$ fixed). Operating thresholds are tuned per class on validation by maximising F_1 over 61 thresholds from 0.05 to 0.95.

5.3. Inference blending

Per-class probabilities are computed as

$$\mathbf{p}(\mathbf{x}) = w^* \sigma(\ell_{\text{fused}}) + (1 - w^*) \sigma(\ell_{\text{ecg}}), \quad (14)$$

where $w^*=1.0$ is fixed *a priori* according to the pre-declared analysis plan. Test predictions are therefore simply the fused-branch outputs for all runs. No blend-weight grid search is performed; the fused head is used directly.

5.4. Ablation variants

We train three variants that differ *only* in the metadata input:

none $\mathbf{x}_{\text{meta}} \leftarrow \mathbf{0}$, $\mathbf{m} \leftarrow \mathbf{0}$; $q_{\text{meta}}=0$ and the metadata embedding is nulled after quality scaling.
demo age and sex are retained; height, weight, BMI values set to 0, their missingness indicators set to 1, and all anthropometric/missingness masks set to 0.
demo+anthro full 8-dim $\mathbf{x}_{\text{meta}}$ as in Eq. (1).

5.5. Seeds, runs, compute

The full training campaign is organised in **six experimental groups (A–F)** totalling **250 unique runs** (270 seed-variant descriptors, of which 20 are shared between groups):

Group A (primary confirmatory, 60 runs). Three metadata variants (none, demo, demo+anthro) $\times$ **20 seeds** $\mathcal{S} = \{2024, \dots, 2043\}$, test fold 10, val fold 9. Python, NumPy, and PyTorch are each re-seeded at the start of every run. Wall-clock time totalled 48.2 h (median 43.8 min/run; range [43.1, 303.5] min, the high value reflecting a transient CPU throttle at epoch 5 with no effect on convergence).

Group B (exploratory, 30 unique runs). Three non-reference meta_hid values {32, 64, 256} $\times$ 10 seeds, demo+anthro condition; reference meta_hid=128 shares seeds with Group A. Note: meta_hid controls both projection and gate hidden dimensions (entangled).

Group C (exploratory, 40 unique runs). Four non-reference λ_{LAUC} values {0.00, 0.04, 0.12, 0.16} $\times$ 10 seeds; reference $\lambda=0.08$ shares seeds with Group A. Note: $w_{\text{fused}} = \max(0, 0.60 - \lambda_{\text{LAUC}})$, so both weights change together.

Group D (exploratory, 20 runs). Augmentation ablation: standard (aug, reference) vs. no-augmentation (noaug), 10 seeds each.

Group E (exploratory, 80 runs). Eight architectural ablations (E1–E8) $\times$ 10 seeds: meta-only fusion (E1), concatenation without gate (E2), additive injection without gate (E3), ReLU gate (E4), no residual injection (E5), availability signal disabled / $q=1$ always (E6), no feature dropout (E7), zero-metadata training / trainmask (E8).

Group F (confirmatory generalisation check, 20 runs). Four alternative fold pairs (test/val: 2/3, 3/4, 7/8, 8/9) $\times$ 5 seeds, demo+anthro condition. No statistical test across the 20 runs as if independent (pre-declared); analysis unit is the fold.

The choice of 20 seeds for Group A is a principled compromise between statistical resolution and the CPU-only computational budget. The paired Wilcoxon test at $n=20$ provides a minimum attainable exact p -value of approximately 0.000002 ($= 2/2^{20}$), enabling detection of small effects such as the demographic-only gain of $\Delta = +0.0007$ macro-AUC. Twenty seeds is above the modal practice in supervised deep learning evaluation, where three to five seeds are commonly reported [37,38], and provides adequate power for the effect sizes observed here ($d_z=0.64$ for demo–none; $d_z=1.65$ for demo+anthro–none).

5.6. Statistical analysis

We report macro-AUC and macro- F_1^* as mean $\pm$ standard deviation across seeds. Marginal 95% percentile bootstrap CIs are obtained by resampling the 20 seed-level values 10,000 times; paired-difference CIs are obtained by resampling the 20 seed-matched differences. Pairwise inference uses exact two-sided paired Wilcoxon signed-rank tests [39] on seed-matched variant differences. We report Cohen's d_z as the main descriptive paired effect size and Hedges-corrected g_z as a small-sample companion.

BH-FDR family (Group A, confirmatory). Benjamini–Hochberg FDR control at $q=0.05$ is applied across the **pre-specified 3-test family**: the three pairwise macro-AUC contrasts in Group A (none vs. demo; none vs. demo+anthro; demo vs. demo+anthro). With $n=20$, the attainable raw two-sided p -values are discrete; the minimum attainable exact Wilcoxon p at $n=20$ is approximately $0.000002 (= 2/2^{20} \approx 1.9 \times 10^{-6})$, achieved when all 20 signed-rank sums equal 0 or 210).

Exploratory groups (B, C, D, E). Results for Groups B, C, D, and E are exploratory. Wilcoxon p -values are reported for descriptive purposes without multiple-comparison correction; no confirmatory inference should be drawn from these values.

Secondary per-class sub-family (DA–NONE, Table 7). As a secondary, post-hoc analysis, BH-FDR correction is additionally applied within the 5-test sub-family of per-class DA–NONE Wilcoxon tests (NORM, MI, STTC, CD, HYP). This sub-family was not pre-specified and serves descriptive purposes only. Per-class tests for the demo–none contrast are exploratory and reported as raw p -values only.

Group F (confirmatory generalisation check). No statistical test across the 20 Group F runs as if independent. Analysis is at the fold level (4 units); the primary quantity is the inter-fold macro-AUC range.

Decision thresholds are tuned per class on fold 9 and fixed for fold 10; reported p -values are therefore conditional on this validation-based model-selection protocol.

5.7. Methodological contribution

The seed-matched paired-Wilcoxon measurement pipeline introduced here is a generalisable framework for multimodal-ML ablations. For each seed s and variant V_k , one trains to obtain a scalar evaluation metric $r_{k,s}$. Pairwise incremental effects are estimated as $\delta_{jk,s} = r_{j,s} - r_{k,s}$. Inference proceeds by applying the exact two-sided Wilcoxon signed-rank test, BH-FDR correction over the pre-specified confirmatory test family, and percentile bootstrap CIs. Effect sizes are reported as Cohen's d_z with Hedges' small-sample correction g_z . Exploratory groups report raw p -values descriptively without FDR correction. The pipeline applies to any setting satisfying three conditions: (i) a baseline model can be paired with an augmented variant across the same random seeds; (ii) the number of seeds is small (say, $n \in [5, 20]$); and (iii) one seeks calibrated inference about the incremental contribution of an additional input modality or component.

6. Results

6.1. Aggregate performance

All reported p -values are conditional on the validation-based model-selection protocol. Table 6 reports macro-AUC and macro- F_1^* for each variant as mean $\pm$ standard deviation across 20 seeds, together with 95% bootstrap CIs and paired Wilcoxon p -values. The demo+anthro variant reaches 0.9289 ± 0.0013 macro-AUC on fold-10, numerically close to the 10-fold CV mean of the xresnet1d101 benchmark of [6] (0.928 ± 0.005 ; direct comparison is limited as our evaluation uses fold-10 alone, which Group F shows is harder than average).

A key finding is that **both demographic-only and full-metadata gains are statistically significant** across 20 paired seeds. The demo–none difference is $+0.0007 \pm 0.0002$ macro-AUC with a 95% paired bootstrap CI of $[0.0002, 0.0011]$ ($p_{\text{raw}}=0.0094$, $p_{\text{BH}}\approx 0.009$, $d_z=0.64$, $g_z=0.61$; 16 of 20 paired seed differences positive). The demo+anthro–none contrast yields $+0.0019 \pm 0.0003$ macro-AUC with a 95% paired bootstrap CI of $[0.0014, 0.0023]$ ($p_{\text{raw}}\approx 0.00001$, $p_{\text{BH}}<0.001$, $d_z=1.65$, $g_z=1.58$; 19 of 20 differences positive). The demographic gain is real but smaller than the anthropometric component; the incremental anthropometric step (demo+anthro–demo) is itself significant ($+0.0012$, $p_{\text{BH}}\approx 0.001$, $d_z=0.95$). These results should be read as protocol-specific measurements rather than as large clinical effect sizes.

6.2. Per-class decomposition

Table 7 and Figs. 3 and 5 report class-wise AUC and the corresponding Δ -AUC for the demo+anthro minus none contrast. With 20 seeds and BH-FDR correction, **three classes reach significance: NORM** ($+0.0016$, $p_{\text{BH}}<0.0001$), **MI** ($+0.0058$, $p_{\text{BH}}<0.0001$), and **STTC** ($+0.0018$, $p_{\text{BH}}=0.00035$). The MI effect has the largest absolute delta ($+0.0058$) and a strong $d_z=2.26$. CD shows a negligible shift (-0.0003 , $p_{\text{BH}}=0.648$). **HYP shows no statistically significant gain** ($+0.0004$, $p_{\text{BH}}=0.648$); we return to this HYP pattern in Section 7.

For the demo–none contrast, exploratory per-class Wilcoxon tests identify significance for NORM ($+0.0006$, $p_{\text{raw}}=0.006$) and MI ($+0.0020$, $p_{\text{raw}}=0.003$), indicating that even demographic variables alone carry a per-class signal in these two morphology-rich classes. These per-class tests are not in the pre-specified BH-FDR family and are reported descriptively.

6.3. Training dynamics

Figure 2 shows the validation macro-AUC trajectory across 10 training epochs for all three variants, displayed as mean $\pm$ standard deviation across all 20 seeds. The cross-seed SD narrows from 0.008–0.010 at epoch 1 to ≈ 0.001 at epoch 10. All three variants converge to near-final performance by epoch 7–8. The final-epoch means and SDs are: none 0.9320 ± 0.0012 , demo 0.9323 ± 0.0008 , demo+anthro 0.9330 ± 0.0012 .

6.4. Missingness robustness

Figure 4 shows inference-time macro-AUC under synthetic anthropometric masking for all three variants across **20 seeds**. At each masking rate $\rho \in \{0, 0.25, 0.5, 0.75, 1.0\}$, each anthropometric field (height, weight, BMI) is independently zeroed with probability ρ at inference time and the availability mask is updated consistently. The stress-test curve remains stable: full anthropometric masking costs only ≈ 0.0014 macro-AUC relative to the unmasked model (Panel A, Δ annotation). At $\rho=1.0$ the demo+anthro model converges to the demo baseline (difference -0.00018 , less than $\frac{1}{6}$ cross-seed SD) and remains $+0.00047$ above the none baseline. Panel B clarifies that the none variant is immune by construction (flat), while demo shows minimal drift (< 0.0001). This demonstrates graceful degradation under deployment-time field dropout.

6.5. Seed distribution and fused-branch gain

Figure 6 shows the distribution of test macro-AUC across all 20 seeds for each variant. The demo+anthro violins are shifted upward relative to none and demo, with minimal within-variant variance ($\text{SD} \leq 0.0013$). The significance brackets confirm the two BH-FDR confirmatory contrasts (Section 5). Across seeds, the fused head consistently outperforms the ECG-only auxiliary head for demo+anthro: the $\text{AUC}_{\text{fused}} - \text{AUC}_{\text{ecg}}$ gap is positive for **all 20 of 20 seeds** (median $+0.0016$; Panel B of Fig. 6), whereas for none and demo this gap is centred near zero (8/20 and 16/20 positive; medians -0.0001 and $+0.0003$

respectively). This is consistent with $w^*=1.0$ fixed a priori: when metadata help, the fused branch captures the gain.

6.6. Comparison with prior work

Tables 8–9 position EZNX-ATLAS-A through a comparability audit rather than within a heterogeneous leaderboard. Our direct numerical comparator is the canonical PTB-XL super-diagnostic benchmark of Strodthoff *et al.* [6]: the none variant lies approximately 0.001 macro-AUC below the 10-fold CV mean (0.9271 vs. 0.928; comparison is indicative as our evaluation uses fold-10 alone, which Group F shows is harder than average) while using a substantially smaller trainable model. The self-supervised and multimodal papers listed in Table 9 address neighbouring questions, but their task definitions, datasets, or metadata modalities differ from our seed-matched PTB-XL demographic/anthropometric ablation.

6.7. Architectural diagnostics and sensitivity ablations

With $w^*=1.0$ fixed a priori, test predictions are those of the fused head alone for all runs.

Groups B and C: hyperparameter sensitivity (exploratory).

Table 1 reports multi-seed (10 seeds) exploratory sensitivity results for Groups B and C. Group B tests three non-reference `meta_hid` values {32, 64, 256} against the reference (128, shared with Group A); all Wilcoxon $p > 0.32$ versus the reference, spanning the macro-AUC range 0.9285–0.9296 with overlapping intervals. Note that `meta_hid` controls both the projection dimension and the gate hidden dimension simultaneously (entangled); varying it does not isolate a single architectural degree of freedom. Group C tests four non-reference λ_{LAUC} values {0.00, 0.04, 0.12, 0.16} against the reference (0.08, shared with Group A); all variants are within the noise of the reference, with p -values ranging from 0.557 to 0.695. Note that $w_{\text{fused}} = \max(0, 0.60 - \lambda_{\text{LAUC}})$, so the two loss weights change simultaneously when λ_{LAUC} is varied. These results confirm architectural robustness rather than identifying an optimal configuration; no BH-FDR correction is applied to exploratory groups.

Table 1. Exploratory sensitivity: fusion capacity (Group B, `meta_hid` controls projection and gate hidden dim jointly) and LAUC-BCE balance (Group C; $w_{\text{fused}} = \max(0, 0.60 - \lambda_{\text{LAUC}})$, both weights change simultaneously). `demo+anthro` condition, 10 seeds. † = default value. Wilcoxon p -values ($n=10$, exact two-sided) are reported descriptively; no FDR correction is applied to exploratory groups.

Group	Parameter / Value	Macro-AUC (mean±SD)	Wilcoxon p vs. ref
B	<code>meta_hid</code> = 32	0.9285 ± 0.0013	0.322
B	<code>meta_hid</code> = 64	0.9285 ± 0.0014	0.432
B	<code>meta_hid</code> = 128† (ref)	0.9289 ± 0.0012	—
B	<code>meta_hid</code> = 256	0.9296 ± 0.0010	0.322
C	λ_{LAUC} = 0.00	0.9286 ± 0.0017	0.557
C	λ_{LAUC} = 0.04	0.9291 ± 0.0012	0.695
C	λ_{LAUC} = 0.08† (ref)	0.9289 ± 0.0012	—
C	λ_{LAUC} = 0.12	0.9288 ± 0.0008	0.625
C	λ_{LAUC} = 0.16	0.9291 ± 0.0008	0.557

Group D: augmentation ablation (exploratory).

Training without stochastic augmentation (`noaug`) yields macro-AUC 0.9275 ± 0.0012 versus 0.9289 ± 0.0012 for the augmented reference ($\Delta = -0.0014$, $p = 0.0195$, $r = 0.725$, $n = 10$). This is the only exploratory group to show a nominally significant difference from reference.

The result suggests that the light augmentation block (Gaussian noise, temporal shift, amplitude jitter) contributes a positive effect and is not a neutral design choice; confirmatory inference would require a larger seed budget.

Group E: architectural ablations (exploratory).

Table 2 reports macro-AUC for eight architectural ablations against the standard configuration (10 seeds each). Two extreme variants show catastrophic degradation: **E1 meta-only** (0.6966 ± 0.0021 , $p=0.002$) confirms that the ECG signal is indispensable, and **E8 trainmask** (0.5046 ± 0.0180 , $p=0.002$)—in which the model is trained with zeroed metadata and then evaluated with real metadata—confirms that metadata must be present during training for the model to learn to use it. All six remaining ablations (E2–E7) show macro-AUC within ± 0.0004 of the reference with $p > 0.27$, indicating that the specific choice of fusion mechanism (concatenation, additive injection, GLU gate) and the inclusion of residual injection, availability signal, and feature dropout do not materially affect macro-AUC.

Table 2. Architecture ablations (Group E, exploratory). *demo+anthro* variant, 10 seeds. E8 (*trainmask*): standard architecture trained with $x_{\text{meta}}=0$, $m=0$; evaluated with real metadata. Wilcoxon p -values ($n=10$, exact two-sided) vs. Standard reference.

Configuration	Macro-AUC (mean±SD)	Wilcoxon p
Standard (reference, Group A)	0.9289 ± 0.0012	—
E1: Meta-only fusion	0.6966 ± 0.0021	0.002
E2: Concatenation (no gate)	0.9290 ± 0.0010	0.922
E3: Additive injection (no gate)	0.9290 ± 0.0012	0.846
E4: ReLU gate (no sigmoid)	0.9288 ± 0.0014	0.770
E5: No residual injection	0.9292 ± 0.0012	0.695
E6: No availability signal ($q=1$)	0.9287 ± 0.0013	0.557
E7: No feature dropout	0.9293 ± 0.0008	0.275
E8: Trainmask (zero meta at training)	0.5046 ± 0.0180	0.002

Group F: multi-split generalisation check (confirmatory).

Table 3 reports macro-AUC for the *demo+anthro* model on four alternative test folds (5 seeds each). The inter-fold range is $[0.9400, 0.9445]$, with a mean of 0.9425 across the four folds. **Fold 10 (Group A reference, 0.9289 ± 0.0013) lies below this range by 0.0111–0.0156 macro-AUC**, confirming that fold 10 is a more challenging evaluation split. Group A results are therefore conservatively estimated and likely to understate the model’s typical performance.

Table 3. Multi-split sensitivity analysis (Group F, confirmatory generalisation check). *demo+anthro* model, 5 seeds per split. Analysis unit: fold (4 units). No statistical test across the 20 runs as if independent (pre-declared). Group A (fold 10) reference included for comparison.

Test fold	Val fold	n seeds	Macro-AUC (mean±SD)
2	3	5	0.9422 ± 0.0007
3	4	5	0.9400 ± 0.0006
7	8	5	0.9434 ± 0.0007
8	9	5	0.9445 ± 0.0013
10 (Group A)	9	20	0.9289 ± 0.0013 (reference)
Inter-fold range (folds 2,3,7,8): $[0.9400, 0.9445]$; mean: 0.9425			

7. Discussion

7.1. Q1 — when do metadata help?

Our 20-seed protocol yields a clear and nuanced answer. **Both demographic and anthropometric variables show statistically significant gains.** Demographics alone (age, sex) yield $\Delta = +0.0007$ macro-AUC ($p_{\text{BH}} \approx 0.009$, $d_z = 0.64$, $n_{\text{pos}} = 16/20$ seeds). This small but reproducible effect is consistent with the partial recoverability argument of Attia *et al.* [17]: age and sex are mostly encoded in the waveform (sex AUC 0.97, age MAE 6.9 y), yet a residual increment not fully captured by the ECG-only model remains detectable with 20 paired seeds. The incremental anthropometric gain beyond demographics is larger and more significant ($\Delta = +0.0012$, $p_{\text{BH}} \approx 0.001$, $d_z = 0.95$): body-size variables provide complementary information not easily recoverable from the 12-lead waveform alone. Both effects remain small in absolute terms ($\Delta < 0.002$ macro-AUC) and should be read as protocol-specific measurements rather than as clinically large effect sizes.

7.2. Q2 — which class benefits from which metadata? The MI/STTC/NORM finding

With 20 seeds and BH-FDR correction, the class-wise effect of full metadata localises to three classes. **MI** shows the strongest effect ($+0.0058$, $p_{\text{BH}} < 0.0001$, $d_z = 2.26$), followed by **STTC** ($+0.0018$, $p_{\text{BH}} = 0.00035$) and **NORM** ($+0.0016$, $p_{\text{BH}} < 0.0001$). Exploratory per-class tests for the demographic-only contrast (demo—none) also identify nominal significance for NORM ($p_{\text{raw}} = 0.006$) and MI ($p_{\text{raw}} = 0.003$); these tests are not in any pre-specified BH-FDR family and are reported descriptively. This pattern suggests that age and sex supply a useful prior for morphology-sensitive classes even when anthropometrics are absent. Plausible mechanisms include body-size-conditioned calibration of ST-segment amplitudes, Q-wave depth, and related markers for MI, and thoracic electrical property effects for STTC.

7.3. The HYP paradox: no detectable gain at the superclass level

Even with 20 seeds, **HYP shows no statistically significant gain** ($+0.0004$, $p_{\text{BH}} = 0.648$). This pattern is consistent with four explanatory hypotheses: (1) low HYP prevalence limits paired test power for a genuinely small effect; (2) the TSPool's max / min voltage statistics already partially absorb body-habitus effects; (3) superclass heterogeneity dilutes a more specific LVH-related signal; and (4) HYP label noise in borderline cases attenuates any body-habitus signal. The non-significant HYP result with 20 seeds strengthens confidence that this is not merely a power limitation, but rather a genuine absence of a detectable effect at the superclass level under this protocol. Exploratory examination of the demo—none HYP test ($\Delta = +0.0015$, $p_{\text{raw}} = 0.048$) reveals a borderline positive demographic signal that does not pass the exploratory threshold of $p < 0.01$ and is not sustained in the full-metadata DA—NONE contrast ($p_{\text{BH}} = 0.648$), consistent with the interpretation that anthropometric variables do not amplify the marginal demographic effect for HYP.

7.4. CD: demographics and the confounding role of age

The demo variant produces a slight negative trend in CD AUC (-0.0011 , $p_{\text{raw}} = 0.123$, exploratory), but this does not survive BH-FDR correction at 20 seeds. This remains consistent with the age-as-confounder hypothesis: directly supplying age may shift the CD decision boundary, but the effect is not reproducibly negative. The demo+anthro CD delta is also non-significant (-0.0003 , $p_{\text{BH}} = 0.648$).

7.5. Q3 — graceful degradation

The joint mechanism of q_{meta} scaling (Eq. (6)), zero-initialised residual injection (Eq. (8)), and sample meta-dropout during training yields a model whose macro-AUC under full anthropometric suppression remains within ≈ 0.0014 of the full unmasked variant

(Fig. 4), converging to the demographics-only baseline. This robustness is demonstrated under synthetic inference-time masking; generalisability to structurally missing data under MCAR, MAR, or MNAR mechanisms remains an open evaluation direction.

7.6. Group D: augmentation is not a neutral choice

The exploratory Group D result ($\Delta = -0.0014$, $p = 0.0195$) reveals that the light stochastic augmentation block—Gaussian noise, temporal shift, and amplitude jitter—contributes a measurable positive effect. While 10 seeds do not enable confirmatory inference, the nominally significant result (the only such finding among all exploratory groups) suggests that practitioners should not treat augmentation as an optional add-on in this setting.

7.7. Q4 — comparability of earlier multimodal gains

Our 20-seed protocol reveals within-variant standard deviations of 0.0010–0.0013 macro-AUC, of the same order as the measured between-variant effects. The audited literature does not provide a directly comparable PTB-XL experiment separating age/sex from height/weight/BMI with repeated seeds, paired tests, and missingness stress testing. We therefore position this paper as providing a narrow but statistically honest reproducibility baseline: +0.0007 from demographics and +0.0019 from full metadata, with class-resolved and missingness-audited evidence.

7.8. Q5 — architectural ablations: ECG is essential, gating is flexible

Group E establishes two key facts. First, the ECG signal is indispensable: the meta-only variant (E1) collapses to 0.6966 macro-AUC, far below any acceptable clinical threshold. Second, the specific fusion mechanism does not drive macro-AUC differences at this scale: all six mechanism variants (E2–E7) achieve 0.9287–0.9293, within the noise of the standard configuration ($p > 0.27$ in all cases). This is notable because E2 (simple concatenation) and E3 (additive injection) achieve performance comparable to the GLU gate, suggesting that the gate's added complexity is not necessary to capture the metadata signal at this effect magnitude. The E8 result (0.5046 macro-AUC) is a negative control confirming that training without metadata prevents the model from learning to use it: the metadata pathway must be active at training time.

7.9. Q6 — generalisation: fold 10 is a conservatively difficult split

Group F shows that the four alternative fold pairs (folds 2, 3, 7, 8 as test folds) all achieve macro-AUC in the range [0.9400, 0.9445], substantially above fold 10 (0.9289). This 0.011–0.016 gap is consistent with fold-level difficulty variation in stratified k-fold settings and confirms that the confirmatory Group A results, conditioned on fold 10, likely *understate* the typical performance of EZNX-ATLAS-A. The conservative nature of fold 10 should be noted when positioning our results relative to benchmarks evaluated on averaged or best-of-fold metrics.

7.10. Sensitivity ablations: architectural robustness confirmed

The multi-seed sensitivity analyses (Groups B and C, Section 6.7) confirm that the reported results are not artefacts of a specific metadata hidden size or AUC-surrogate weight. All seven non-reference variants (three metaH, four lauc) show Wilcoxon $p > 0.32$ against the standard configuration, with macro-AUC ranges spanning less than 0.0012, providing statistically grounded evidence of architectural robustness.

7.11. Why small measured effects remain informative

Reporting +0.0007 and +0.0019 macro-AUC gains is not a failure; it is a refinement of the multimodal claim. The paired seed-level 95% CIs are strictly positive for the full-

metadata contrast ($[0.0014, 0.0023]$), supporting a reproducible directional effect without implying a large clinical effect size. The demographic-only CI ($[0.0002, 0.0011]$) is positive but narrow, consistent with a small real effect that required 20 seeds to detect. These results indicate that: (i) a strong ECG-only model is already close to the incremental information supplied by the tested metadata; (ii) the residual signal localises to morphology-rich classes (MI, STTC, NORM); (iii) demographic fusion alone can create class-specific shifts auditable only through paired multi-seed statistics; and (iv) anthropometrics contribute a second, statistically significant increment beyond demographics.

7.12. Limitations

1. **Evaluation scope and fold difficulty.** The confirmatory Group A results are conditioned on fold 10, which Group F identifies as lying below the inter-fold macro-AUC range $[0.9400, 0.9445]$ by approximately 0.011–0.016. This indicates that fold 10 is a more challenging split and that the reported gains are conservative estimates. No external cohort evaluation has been performed, limiting generalisation across institutions, acquisition devices, demographic distributions, and temporal shifts.
2. **Task and metadata scope.** The analysis is restricted to the 5-superclass task and to age, sex, height, weight, and BMI. Variables such as blood pressure, electrolytes, and echocardiographic measurements are absent. Age, sex, and body size are not purely biological descriptors—they also carry population-structure signal—so measured gains cannot be unambiguously attributed to physiological mechanisms alone.
3. **Model selection and downstream metrics.** Class-wise F_1 thresholds and checkpoint selection are all determined on fold 9 and fixed for fold 10. The reported p -values are therefore conditional on this validation-based selection protocol, and F_1 should be read more cautiously than AUC.
4. **Architectural attribution.** Although Group E ablations (E2–E7) show that individual components do not materially affect macro-AUC, the entangled nature of certain hyperparameters (e.g., `meta_hid` in Group B controls both projection and gate dimensions) prevents clean attribution. The metadata MLPs retain BatchNorm, which may be less stable than layer-normalised alternatives when metadata-rich samples are rare within a mini-batch.
5. **Statistical granularity.** The reported bootstrap intervals quantify variability across seed-level runs rather than patient-level resampling, and the seed-matched comparisons share the same dataset, preprocessing pipeline, and split definition. The paired contrasts are therefore correlated repeated runs on one fixed cohort split rather than independent cohort replications. Exploratory group results (B, C, D, E) report raw p -values without FDR correction and should not support confirmatory conclusions.
6. **Reproducibility provenance.** All 250 runs were executed on a CPU-only machine (Intel Core i5, 8 GB RAM) using PyTorch 2.3.1+cpu in a local Jupyter Notebook environment under Windows. CPU-only execution eliminates CUDA non-determinism from parallel atomic operations, providing a stronger reproducibility guarantee than GPU training for the same seeds on identical hardware and software; results may nonetheless differ across different CPU architectures or BLAS library versions.

8. Conclusions

EZNX-ATLAS-A reaches macro-AUC comparable to the canonical PTB-XL superclass benchmark while serving as an instrument for measuring the incremental contribution of clinical metadata across a **250-run campaign organised in six experimental groups (A–F)**. The main findings are:

- (a) **Demographic variables** (age, sex) **add a small but statistically significant gain** ($\Delta = +0.0007$ macro-AUC, 95% paired seed-level CI [0.0002, 0.0011], $p_{BH} \approx 0.009$); the effect localises to NORM and MI and is consistent with a residual demographic signal not fully recovered from the waveform alone.
- (b) **Anthropometric variables** (height, weight, BMI) **provide an additional highly significant gain** of +0.0012 over demographics ($p_{BH} \approx 0.001$), for a total full-metadata gain of +0.0019 over ECG alone ($p_{BH} < 0.001$, 95% CI [0.0014, 0.0023]); the class-wise signal concentrates in MI, STTC, and NORM.
- (c) **HYP shows no significant effect** under this superclass protocol ($p_{BH} = 0.648$), consistent with voltage-based LVH criteria being diluted by superclass heterogeneity, and **CD shows no significant effect** ($p_{BH} = 0.648$), despite the demographic CD hypothesis.
- (d) **Quality-gated fusion** with zero-initialised residual injection remains stable under simulated missingness, losing only ≈ 0.0014 macro-AUC under full anthropometric masking, converging to the demographics-only baseline.
- (e) **Exploratory Groups B and C** (hyperparameter sensitivity, 70 unique runs) confirm architectural robustness: metadata hidden size and AUC-surrogate weight show no statistically significant effect on macro-AUC (all Wilcoxon $p > 0.32$, 10 seeds each).
- (f) **Group D** (augmentation ablation, 20 runs) finds that removing light stochastic augmentation causes a nominally significant macro-AUC drop ($\Delta = -0.0014$, $p = 0.0195$), indicating that augmentation is not a neutral design choice.
- (g) **Group E** (architectural ablations, 80 runs) confirms that the ECG signal is indispensable (meta-only AUC 0.6966) and that metadata must be present during training (E8 trainmask AUC 0.5046); the specific fusion mechanism does not materially affect macro-AUC.
- (h) **Group F** (multi-split check, 20 runs) shows that fold 10 lies below the inter-fold range [0.9400, 0.9445] from four alternative splits, confirming that Group A results are conservatively estimated.

More broadly, the seed-matched paired-Wilcoxon + BH-FDR (over the pre-declared 3-test confirmatory family) + paired-bootstrap pipeline described in §5 constitutes a reusable methodological baseline for any multimodal-ML ablation in which an additional input modality is evaluated across repeated runs. The 20-seed design for Group A provides enough paired runs to detect small effects such as the demographic gain ($d_z = 0.64$), highlighting the importance of adequate seed budgets in multimodal deep learning evaluation.

Supplementary Materials: Supplementary materials bundled with the release snapshot include: **Table S1** — the 3-test BH-FDR pre-specified family (Group A pairwise macro-AUC contrasts) with raw and adjusted p -values; seed-level JSON outputs for all 250 runs across six groups; seed-level CSV/Markdown statistical tables; dataset-integrity and statistical-protocol notes; environment files; and source/figure-generation scripts.

Author Contributions: Conceptualization, E.S., M.M., and B.M.A.; methodology, E.S., Y.O., K.E.W., M.C.T., and D.H.P.-O.; software, E.S.; validation, E.S., Y.O. and M.C.T.; formal analysis, E.S., K.E.W. and M.C.T.; investigation, E.S.; resources, E.S.; data curation, E.S.; writing—original draft preparation, E.S.; writing—review and editing, E.S., Y.O., K.E.W., M.C.T., M.M., D.H.P.-O., and B.M.A.; visualization, E.S.; supervision, Y.O., K.E.W., M.C.T., and D.H.P.-O.; project administration, E.S. All authors have read and approved the final version of the manuscript as submitted.

Funding: This research received no external funding.

Institutional Review Board Statement: Ethical review and institutional approval were waived for this study because it uses only the public, fully de-identified PTB-XL dataset [5], whose original collection was approved by the institutional ethics committee of PTB Berlin.

Informed Consent Statement: Patient consent was waived because this work uses only the public, fully de-identified PTB-XL dataset, for which informed consent was obtained at initial collection according to the PTB Berlin protocol.

Data Availability Statement: The PTB-XL v1.0.3 dataset is freely available on PhysioNet (<https://physionet.org/content/ptb-xl/1.0.3/>, doi:10.13026/kfzx-aw45) [5,36]. The complete code, all seed-level result JSON files for all 250 runs across six experimental groups, derived statistical artifacts, and the six manuscript figures are publicly available at https://github.com/ezyneegnane/ezyx-atlas-a_github_code (MIT licence). Patient-level probability arrays for all 60 Group A runs were not archived in the release and are therefore not included in the package.

Acknowledgments: The authors thank the PTB-XL authors for releasing the dataset and the Strodtzoff group for the open benchmark code whose experimental spirit this protocol extends.

Use of Generative AI: During the preparation of this submission, the authors used Claude (Anthropic) to audit the reference list for dataset and terminology consistency, to verify manuscript compliance with journal formatting requirements, and to correct grammar, spelling, and punctuation errors in the manuscript text. The authors reviewed all AI-generated suggestions and take full responsibility for the content of this publication.

Conflicts of Interest: No conflict of interest is declared by any of the authors.

Abbreviations used in this manuscript:

AUC	Area under the ROC curve
BCE	Binary cross-entropy
BH	Benjamini–Hochberg
BMI	Body mass index
CD	Conduction disturbance (PTB-XL superclass)
ECG	Electrocardiogram
EHR	Electronic health record
FDR	False discovery rate
GLU	Gated Linear Unit
HYP	Hypertrophy (PTB-XL superclass)
IQR	Interquartile range
LVH	Left-ventricular hypertrophy
MAR	Missing at random
MCAR	Missing completely at random
MI	Myocardial infarction (PTB-XL superclass)
MLP	Multilayer perceptron
MNAR	Missing not at random
NORM	Normal ECG (PTB-XL superclass)
PTB-XL	Physikalisch-Technische Bundesanstalt eXtra Large ECG corpus
SCP	Standard Communications Protocol (for ECG coding, ISO 11073-91064)
STTC	ST/T change (PTB-XL superclass)
SSL	Self-supervised learning
Superclass	One of the five PTB-XL top-level diagnostic categories (NORM, MI, STTC, CD, HYP)

Table 4. PTB-XL 1.0.3 dataset statistics used in this study. Superclass prevalence is the fraction of records carrying that superclass (multi-label; rows do not sum to 100%). Positive weights are derived via Eq. (2) from training-fold prevalences.

	NORM	MI	STTC	CD	HYP
Training prevalence (f_c , %)	43.6	25.1	24.0	22.4	12.2
Positive weight w_c^+	1.29	2.98	3.16	3.46	7.22
Validation prevalence (%)	43.7	24.7	24.2	22.7	12.3
Test prevalence (%)	43.8	25.0	23.7	22.6	11.9

Training / validation / test records (folds 1–8 / 9 / 10)	17,418 / 2,183 / 2,198
Metadata completeness (age / sex / height / weight / BMI)	98.7 / 100.0 / 31.9 / 43.2 / 30.9 %

Table 5. Hyperparameters and compute cost. All values are held constant across the 60 Group A runs (20 seeds $\times$ 3 variants) and the exploratory/confirmatory groups. All runs were executed on a CPU-only machine (Intel Core i5, 8 GB RAM, PyTorch 2.3.1+cpu) with no GPU or CUDA device.

Hyperparameter	Value
Optimizer	AdamW
Learning rate	10^{-3}
Weight decay	5×10^{-4}
LR schedule	CosineAnnealingWarmRestarts, $T_0=10$, $T_{\text{mult}}=2$
Batch size (effective)	$32 \times 2 = 64$
Gradient clipping	maximum norm 1.0
Epochs	10
Loss weights (fused / ECG / meta / AUC surrogate)	0.52 / 0.30 / 0.10 / 0.08 (reference; Group C varies λ_{LAUC} ; effective total <1.00 when $q_{\text{meta}} < 1$)
Sample meta-dropout p	0.10 (training only)
Inference blend w^*	Fixed at 1.00 (a priori, all runs)
Parameters (trainable)	3.95 M
Training-device provenance	Intel Core i5 CPU, 8 GB RAM; PyTorch 2.3.1+cpu; CPU-only, no CUDA/GPU
Determinism flags	Code includes CuDNN deterministic/benchmark flags (inert on CPU-only); CPU execution eliminates CUDA non-determinism.
Wall-clock time per run (Group A)	43.8 min median (48.2 min mean; 43.1 to 303.5 min range)
Total compute (Group A, 60 runs)	48.2 h
Total compute (all 250 runs)	≈ 185 h

Table 6. Primary confirmatory results (Group A) on the PTB-XL fold-10 test set. Each variant was trained from **20 different random seeds**; macro-AUC and macro- F_1^* are reported as mean $\pm$ SD with 95% percentile bootstrap CIs obtained by resampling the 20 seed indices 10,000 times. For each pairwise contrast: average within-seed difference $\pm$ SEM, 95% bootstrap CI, two-sided exact Wilcoxon signed-rank p , BH-FDR-adjusted p over the **pre-specified 3-test confirmatory family**, Cohen's d_z , and Hedges g_z .

Variant	Macro-AUC	Macro- F_1^*	95% CI / diff. CI
none	0.9271 ± 0.0011	0.7441 ± 0.0030	[0.9266, 0.9276]
demo	0.9277 ± 0.0010	0.7458 ± 0.0028	[0.9272, 0.9282]
demo+anthro	0.9289 ± 0.0013	0.7462 ± 0.0034	[0.9283, 0.9295]
Paired comparisons (macro-AUC)			
demo – none	$+0.0007 \pm 0.0002$	$p_{\text{raw}}=0.0094$, $p_{\text{BH}} \approx 0.009$, $d_z=0.64$, $g_z=0.61$	diff. CI [0.0002, 0.0011]
demo+anthro – demo	$+0.0012 \pm 0.0003$	$p_{\text{raw}}=0.0006$, $p_{\text{BH}} \approx 0.001$, $d_z=0.95$, $g_z=0.91$	diff. CI [0.0007, 0.0018]
demo+anthro – none	$+0.0019 \pm 0.0003$	$p_{\text{raw}} \approx 0.00001$, $p_{\text{BH}} < 0.001$, $d_z=1.65$, $g_z=1.58$	diff. CI [0.0014, 0.0023]

Table 7. Per-class AUC on PTB-XL fold-10 (Group A), mean $\pm$ standard deviation across **20 seeds**, and paired Wilcoxon p -values for $\Delta_{\text{AUC}} = \text{demo+anthro} - \text{none}$. BH-FDR correction is applied within the 5-test per-class DA–NONE sub-family (secondary analysis, not pre-specified); bold entries reach significance. NORM, MI, and STTC reach significance (all $p_{\text{BH}} < 0.001$). CD and HYP share the same BH-adjusted p (0.648) because the BH step-down monotonicity rule reduces HYP’s initial adjustment (0.745) to match CD’s.

Variant	NORM	MI	STTC	CD	HYP
none	0.9481 $\pm$ 0.0013	0.9245 $\pm$ 0.0021	0.9355 $\pm$ 0.0014	0.9179 $\pm$ 0.0021	0.9093 $\pm$ 0.0023
demo	0.9486 $\pm$ 0.0011	0.9265 $\pm$ 0.0025	0.9357 $\pm$ 0.0015	0.9169 $\pm$ 0.0022	0.9108 $\pm$ 0.0024
demo+anthro	0.9497 $\pm$ 0.0013	0.9303 $\pm$ 0.0028	0.9373 $\pm$ 0.0014	0.9177 $\pm$ 0.0019	0.9097 $\pm$ 0.0036
$\Delta_{\text{AUC}} (\text{demo+anthro} - \text{none})$	+0.0016	+0.0058	+0.0018	−0.0003	+0.0004
Wilcoxon p -value (raw)	<0.001	<0.001	0.0002	0.648	0.596
BH-FDR p -value	<0.0001	<0.0001	0.00035	0.648	0.648
Significant after BH-FDR	yes	yes	yes	no	no

Table 8. Direct numerical comparators on PTB-XL. Works evaluated on the PTB-XL super-diagnostic benchmark against which the present results can be directly compared.

Ref.	Audited setting	Metadata / fusion	Use in this manuscript
[6]	PTB-XL super-diagnostic benchmark, official folds	Waveform-only models	ECG-only reference: xresnet1d101 achieves 10-fold CV macro-AUC 0.928 ± 0.005 ; our none variant scores 0.9271 on fold-10 alone. Comparison is indicative: fold-10 is harder than average per the Group F fold-sensitivity analysis.
[12]	PTB-X/PTB-XL arrhythmia classification with 2, 5, 10, and 15 classes	Raw ECG with age/gender branches in CNN/VGG variants	The closest published PTB-XL paper to include demographic branches; however, it reports accuracy-based multi-granularity metrics on a different task structure and provides no anthropometric ablation or missingness test.
[13]	Raw-signal, feature-only, and combined models for PTB-XL multi-label diagnosis	PTB-XL+ tables fused with raw signal	Useful PTB-XL multi-label context; the auxiliary inputs are ECG-derived signal features, not patient-record anthropometrics.
This work	PTB-XL super-diagnostic task, official folds, 20 seeds per variant	none vs. demo vs. full metadata, with availability masks	Demographics yield +0.0007 macro-AUC (localised in NORM and MI) and full metadata yield +0.0019 macro-AUC (concentrated in NORM, MI, and STTC) over ECG alone; both BH-significant. Quality-gated fusion remains robust under anthropometric dropout.

Table 9. Adjacent works whose task, dataset, or metadata definition differs from the present PTB-XL demographic/anthropometric ablation. Listed for context; direct numerical comparison is not appropriate.

Ref.	Audited setting	Metadata / fusion	Use in this manuscript
[21]	PTB-XL+ resource and associated benchmark	Uni-G, 12SL, and ECGDeli tables	Relevant same-corpus context, but the added inputs are derived from the ECG rather than from patient-record variables.
[7–9,19]	Self-supervised or foundation ECG pretraining with PTB-XL evaluations	Waveform-only pretraining / fine-tuning	Self-supervised objectives on raw ECG waveforms; no patient metadata involved.
[10]	AF-recurrence prediction after ablation on a non-PTB-XL cohort	Surface ECG + clinical variables	Useful multimodal example, but outside PTB-XL super-diagnostic classification.
[11]	rECGnition_v2.0 arrhythmia classification on MIT-BIH, INCARTDB, and the European ST-T Database	Patient fusion via self-attentive canonical correlation	Methodological fusion context, without a PTB-XL metadata comparator.




Figure 2. Validation macro-AUC trajectory across 10 training epochs, displayed as mean $\pm$ standard deviation across all **20 seeds** per variant (shaded band = ± 1 SD; evaluated on fold 9). Cross-seed SD narrows from 0.008–0.010 at epoch 1 to ≈ 0.001 at epoch 10. Final-epoch means and SDs: none 0.9320 ± 0.0012 , demo 0.9323 ± 0.0008 , demo+anthro 0.9330 ± 0.0012 . The inter-variant gap remains within one cross-seed SD throughout, consistent with the finding that the $+0.0019$ macro-AUC effect is only detectable through seed-level inference.




Figure 3. Per-class Δ -AUC of demo and demo+anthro relative to none, mean $\pm$ paired SEM across **20 seeds**. **Orange** markers (DA–none): *** = $p_{\text{BH}} < 0.001$ (BH-FDR corrected within the per-class DA–NONE sub-family; all three significant classes NORM, MI, STTC reach $p_{\text{BH}} < 0.001$). **Blue** markers (DEMO–none): * = $p_{\text{raw}} < 0.01$ (exploratory, not in the pre-specified BH family; NORM $p = 0.006$, MI $p = 0.003$). CD and HYP show no significant gain in either contrast.




Figure 4. Missingness robustness of quality-gated fusion across **20 seeds**. **(A)** Test macro-AUC as a function of the fraction ρ of anthropometric fields withheld at inference (MCAR); dashed horizontal lines show each variant's $\rho=0$ baseline; the double-headed arrow quantifies the degradation cost. **(B)** Normalised degradation $\Delta AUC = AUC(\rho) - AUC(0)$ in units of 10^{-3} for direct cross-variant comparison. Full masking ($\rho=1.0$) costs ≈ 0.0014 macro-AUC for demo+anthro, converging to the demo baseline (difference -0.00018 , within $\frac{1}{6}$ SD); the none variant is immune by construction. Availability-mask updating is performed consistently at each rate.




Figure 5. Per-class performance decomposition across **20 seeds**. **(A)** Per-class AUC for the three ablation variants (mean $\pm$ SD); colormap is normalised within each class column so that within-class variant differences are perceptually visible. **(B)** ΔAUC delta matrix: top row = demo—none (exploratory, threshold $p < 0.01$); bottom row = demo+anthro—none (BH-FDR corrected within per-class sub-family). Each cell shows Δ mean and Cohen's d_z . MI shows the largest absolute shift ($+0.0058$, $d_z = 2.26$, $p_{BH} < 0.0001$). NORM and MI already reach significance under demographics alone (top row, $p < 0.01$); STTC additionally reaches significance only with full anthropometrics (bottom row, $p_{BH} = 0.00035$). CD and HYP show no significant gain in either contrast.




Figure 6. Multi-seed reproducibility and architectural validation across **20 seeds**. **(A)** Macro-AUC distribution per variant (violin = kernel density; thick bar = IQR; horizontal tick = median; dashed line = mean; dots = individual seeds; gray lines = paired seed trajectories). All three BH-FDR confirmatory brackets are shown: demo–none ($p_{BH} \approx 0.009$, **), demo+anthro–demo ($p_{BH} \approx 0.001$, ***), demo+anthro–none ($p_{BH} < 0.001$, ***); significance key: ** $p_{BH} < 0.01$; *** $p_{BH} < 0.001$. Gray paired trajectories show that demo+anthro exceeds none in 19 of 20 seeds and demo in 16 of 20. **(B)** $AUC_{fused} - AUC_{ECG}$ per seed: the fused branch exceeds the ECG-only auxiliary head for all 20/20 seeds under demo+anthro (median +0.0016), confirming that the fusion mechanism captures the metadata signal at inference time. The gap is near-zero for none (8/20, median −0.0001) and intermediate for demo (16/20, median +0.0003).

References

- Hannun, A.Y.; Rajpurkar, P.; Haghpanahi, M.; Tison, G.H.; Bourn, C.; Turakhia, M.P.; Ng, A.Y. Cardiologist-level arrhythmia detection and classification in ambulatory electrocardiograms using a deep neural network. *Nature Medicine* **2019**, *25*, 65–69. <https://doi.org/10.1038/s41591-018-0268-3>.
- Ribeiro, A.H.; Ribeiro, M.H.; Paixão, G.M.M.; Oliveira, D.M.; Gomes, P.R.; Canazart, J.A.; Ferreira, M.P.S.; Andersson, C.R.; Macfarlane, P.W.; Meira Jr, W.; et al. Automatic diagnosis of the 12-lead ECG using a deep neural network. *Nature Communications* **2020**, *11*, 1760. <https://doi.org/10.1038/s41467-020-15432-4>.
- Attia, Z.I.; Kapa, S.; Lopez-Jimenez, F.; McKie, P.M.; Ladewig, D.J.; Satam, G.; Pellikka, P.A.; Enriquez-Sarano, M.; Noseworthy, P.A.; Munger, T.M.; et al. Screening for cardiac contractile dysfunction using an artificial intelligence-enabled electrocardiogram. *Nature Medicine* **2019**, *25*, 70–74. <https://doi.org/10.1038/s41591-018-0240-2>.
- Siontis, K.C.; Noseworthy, P.A.; Attia, Z.I.; Friedman, P.A. Artificial intelligence-enhanced electrocardiography in cardiovascular disease management. *Nature Reviews Cardiology* **2021**, *18*, 465–478. <https://doi.org/10.1038/s41569-020-00503-2>.
- Wagner, P.; Strodthoff, N.; Bousseljot, R.D.; Kreiseler, D.; Lunze, F.I.; Samek, W.; Schaeffter, T. PTB-XL, a large publicly available electrocardiography dataset. *Scientific Data* **2020**, *7*, 154. <https://doi.org/10.1038/s41597-020-0495-6>.
- Strodthoff, N.; Wagner, P.; Schaeffter, T.; Samek, W. Deep Learning for ECG Analysis: Benchmarks and Insights from PTB-XL. *IEEE Journal of Biomedical and Health Informatics* **2021**, *25*, 1519–1528. <https://doi.org/10.1109/JBHI.2020.3022989>.
- Mehari, T.; Strodthoff, N. Self-supervised representation learning from 12-lead ECG data. *Computers in Biology and Medicine* **2022**, *141*, 105114. <https://doi.org/10.1016/j.compbiomed.2021.105114>.
- Weimann, K.; Conrad, T.O.F. Self-Supervised Pre-Training with Joint-Embedding Predictive Architecture Boosts ECG Classification Performance. *Computers in Biology and Medicine* **2025**, *196*, 110809. <https://doi.org/10.1016/j.compbiomed.2025.110809>.
- Coppola, E.E.; Savardi, M.; Massucci, M.; Adamo, M.; Metra, M.; Signoroni, A. HuBERT-ECG: a self-supervised foundation model for broad and scalable cardiac applications, 2024. medRxiv preprint (not peer-reviewed), <https://doi.org/10.1101/2024.11.14.24317328>.

10. Qiu, Y.; Guo, H.; Wang, S.; Yang, S.; Peng, X.; Xiayao, D.; Chen, R.; Yang, J.; Liu, J.; Li, M.; et al. Deep learning-based multimodal fusion of the surface ECG and clinical features in prediction of atrial fibrillation recurrence following catheter ablation. *BMC Medical Informatics and Decision Making* **2024**, *24*, 225. <https://doi.org/10.1186/s12911-024-02616-x>.
11. Srivastava, S.; Kumar, D.; Jiwari, R.; Seth, S.; Sharma, D. rECGnition_v2.0: Self-Attentive Canonical Fusion of ECG and Patient Data using Deep Learning for Effective Cardiac Diagnostics, 2025. arXiv preprint arXiv:2502.16255, <https://doi.org/10.48550/arXiv.2502.16255>.
12. Atwa, A.E.M.; Atlam, E.S.; Ahmed, A.; Atwa, M.A.; Abdelrahim, E.M.; Siam, A.I. Interpretable Deep Learning Models for Arrhythmia Classification Based on ECG Signals Using PTB-X Dataset. *Diagnostics* **2025**, *15*, 1950. <https://doi.org/10.3390/diagnostics15151950>.
13. Holloway, C.; Saheer, L.B.; Dar, M.M. Evaluation of Raw Signal and Feature-Based Deep Learning Models for Multi-Label ECG Diagnosis on PTB-XL. In Proceedings of the Artificial Intelligence XLII (SGAI-AI). Springer, 2026, Vol. 16302, LNCS. In press; page numbers not yet assigned.
14. Nonaka, N.; Seita, J. In-depth Benchmarking of Deep Neural Network Architectures for ECG Diagnosis. In Proceedings of the Proceedings of the 6th Machine Learning for Healthcare Conference, 2021, Vol. 149, PMLR, pp. 414–439.
15. Śmigiel, S.; Pałczyński, K.; Ledziński, D. ECG Signal Classification Using Deep Learning Techniques Based on the PTB-XL Dataset. *Entropy* **2021**, *23*, 1121. <https://doi.org/10.3390/e23091121>.
16. Pałczyński, K.; Śmigiel, S.; Ledziński, D.; Bujnowski, S. Study of the Few-Shot Learning for ECG Classification Based on the PTB-XL Dataset. *Sensors* **2022**, *22*, 904. <https://doi.org/10.3390/s22030904>.
17. Attia, Z.I.; Friedman, P.A.; Noseworthy, P.A.; Lopez-Jimenez, F.; Ladewig, D.J.; Satam, G.; Pellikka, P.A.; Munger, T.M.; Asirvatham, S.J.; Scott, C.G.; et al. Age and Sex Estimation Using Artificial Intelligence From Standard 12-Lead ECGs. *Circulation: Arrhythmia and Electrophysiology* **2019**, *12*, e007284. <https://doi.org/10.1161/CIRCEP.119.007284>.
18. Kiyasseh, D.; Zhu, T.; Clifton, D.A. CLOCS: Contrastive Learning of Cardiac Signals Across Space, Time, and Patients. In Proceedings of the Proceedings of the 38th International Conference on Machine Learning, 2021, Vol. 139, PMLR, pp. 5606–5615.
19. Na, Y.; Park, M.; Tae, Y.; Joo, S. Guiding Masked Representation Learning to Capture Spatio-Temporal Relationship of Electrocardiogram. In Proceedings of the International Conference on Learning Representations (ICLR), 2024.
20. Nguyen, H.D.; Pham, T.T.; Le, N.; Nguyen, V. TolerantECG: A Foundation Model for Imperfect Electrocardiogram, 2025. arXiv preprint arXiv:2507.09887, <https://doi.org/10.48550/arXiv.2507.09887>.
21. Strodthoff, N.; Mehari, T.; Nagel, C.; Aston, P.J.; Sundar, A.; Graff, C.; Kanters, J.K.; Haverkamp, W.; Dössel, O.; Loewe, A.; et al. PTB-XL+, a comprehensive electrocardiographic feature dataset. *Scientific Data* **2023**, *10*, 279. <https://doi.org/10.1038/s41597-023-02153-8>.
22. Dauphin, Y.N.; Fan, A.; Auli, M.; Grangier, D. Language Modeling with Gated Convolutional Networks. In Proceedings of the Proceedings of the 34th International Conference on Machine Learning, 2017, Vol. 70, PMLR, pp. 933–941.
23. Arevalo, J.; Solorio, T.; Montes-y Gómez, M.; González, F.A. Gated Multimodal Networks. *Neural Computing and Applications* **2020**, *32*, 10209–10228. <https://doi.org/10.1007/s00521-019-04559-1>.
24. Chen, C.; Dou, Q.; Jin, Y.; Chen, H.; Qin, J.; Heng, P.A. Robust Multimodal Brain Tumor Segmentation via Feature Disentanglement and Gated Fusion. In Proceedings of the MICCAI, 2019, Vol. 11766, LNCS, pp. 447–456. https://doi.org/10.1007/978-3-030-32248-9_50.
25. Kim, J.; Koh, J.; Kim, Y.; Choi, J.; Hwang, Y.; Choi, J.W. Robust Deep Multi-modal Learning Based on Gated Information Fusion Network. In Proceedings of the Asian Conference on Computer Vision (ACCV), 2018. https://doi.org/10.1007/978-3-030-20870-7_6.
26. Baltrušaitis, T.; Ahuja, C.; Morency, L.P. Multimodal Machine Learning: A Survey and Taxonomy. *IEEE Transactions on Pattern Analysis and Machine Intelligence* **2019**, *41*, 423–443. <https://doi.org/10.1109/TPAMI.2018.2798607>.

27. Havaei, M.; Guizard, N.; Chapados, N.; Bengio, Y. HeMIS: Hetero-Modal Image Segmentation. In Proceedings of the MICCAI, 2016, Vol. 9901, *LNCS*, pp. 469–477. https://doi.org/10.1007/978-3-319-46723-8_54.
28. Ma, M.; Ren, J.; Zhao, L.; Tulyakov, S.; Wu, C.; Peng, X. SMIL: Multimodal Learning with Severely Missing Modality. In Proceedings of the Proceedings of the AAAI Conference on Artificial Intelligence, 2021, Vol. 35, pp. 2302–2310. <https://doi.org/10.1609/aaai.v35i3.16330>.
29. Zhang, Y.; He, N.; Yang, J.; Li, Y.; Wei, D.; Huang, Y.; Zhang, Y.; He, Z.; Zheng, Y. mmFormer: Multimodal Medical Transformer for Incomplete Multimodal Learning of Brain Tumor Segmentation. In Proceedings of the MICCAI, 2022, Vol. 13435, *LNCS*, pp. 107–117. https://doi.org/10.1007/978-3-031-16443-9_11.
30. Sokolow, M.; Lyon, T.P. The Ventricular Complex in Left Ventricular Hypertrophy as Obtained by Unipolar Precordial and Limb Leads. *American Heart Journal* **1949**, *37*, 161–186. [https://doi.org/10.1016/0002-8703\(49\)90562-1](https://doi.org/10.1016/0002-8703(49)90562-1).
31. Casale, P.N.; Devereux, R.B.; Kligfield, P.; Eisenberg, R.R.; Miller, D.H.; Chaudhary, B.S.; Phillips, M.C. Electrocardiographic Detection of Left Ventricular Hypertrophy: Development and Prospective Validation of Improved Criteria. *Journal of the American College of Cardiology* **1985**, *6*, 572–580. [https://doi.org/10.1016/S0735-1097\(85\)80115-7](https://doi.org/10.1016/S0735-1097(85)80115-7).
32. Romhilt, D.W.; Estes, E.H. A Point-Score System for the ECG Diagnosis of Left Ventricular Hypertrophy. *American Heart Journal* **1968**, *75*, 752–758. [https://doi.org/10.1016/0002-8703\(68\)90035-5](https://doi.org/10.1016/0002-8703(68)90035-5).
33. Peguero, J.G.; Lo Presti, S.; Perez, J.; Issa, O.; Brenes, J.C.; Tolentino, A. Electrocardiographic Criteria for the Diagnosis of Left Ventricular Hypertrophy. *Journal of the American College of Cardiology* **2017**, *69*, 1694–1703. <https://doi.org/10.1016/j.jacc.2017.01.037>.
34. Rautaharju, P.M.; Surawicz, B.; Gettes, L.S. AHA/ACCF/HRS Recommendations for the Standardization and Interpretation of the Electrocardiogram: Part IV: The ST Segment, T and U Waves, and the QT Interval. *Circulation* **2009**, *119*, e241–e250. <https://doi.org/10.1161/CIRCULATIONAHA.108.191096>.
35. Macfarlane, P.W.; van Oosterom, A.; Pahlm, O.; Kligfield, P.; Janse, M.; Camm, J., Eds. *Comprehensive Electrocardiology*, 2nd ed.; Springer: London, 2011. <https://doi.org/10.1007/978-0-85729-871-3>.
36. Goldberger, A.L.; Amaral, L.A.N.; Glass, L.; Hausdorff, J.M.; Ivanov, P.C.; Mark, R.G.; Mietus, J.E.; Moody, G.B.; Peng, C.K.; Stanley, H.E. PhysioBank, PhysioToolkit, and PhysioNet: Components of a New Research Resource for Complex Physiologic Signals. *Circulation* **2000**, *101*, e215–e220. <https://doi.org/10.1161/01.CIR.101.23.e215>.
37. Bouthillier, X.; Delaunay, P.; Bronzi, M.; Trofimov, A.; Nichyporuk, B.; Szeto, J.; Mohammadi Sepahvand, N.; Raff, E.; Madan, K.; Voleti, V.; et al. Accounting for Variance in Machine Learning Benchmarks. In Proceedings of the Proceedings of Machine Learning and Systems, 2021, Vol. 3, pp. 747–769.
38. Dodge, J.; Gururangan, S.; Card, D.; Schwartz, R.; Smith, N.A. Show Your Work: Improved Reporting of Experimental Results. In Proceedings of the Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing. Association for Computational Linguistics, 2019, pp. 2185–2194. <https://doi.org/10.18653/v1/D19-1224>.
39. Hollander, M.; Wolfe, D.A.; Chicken, E. *Nonparametric Statistical Methods*, 3rd ed.; John Wiley & Sons: Hoboken, NJ, 2014.